



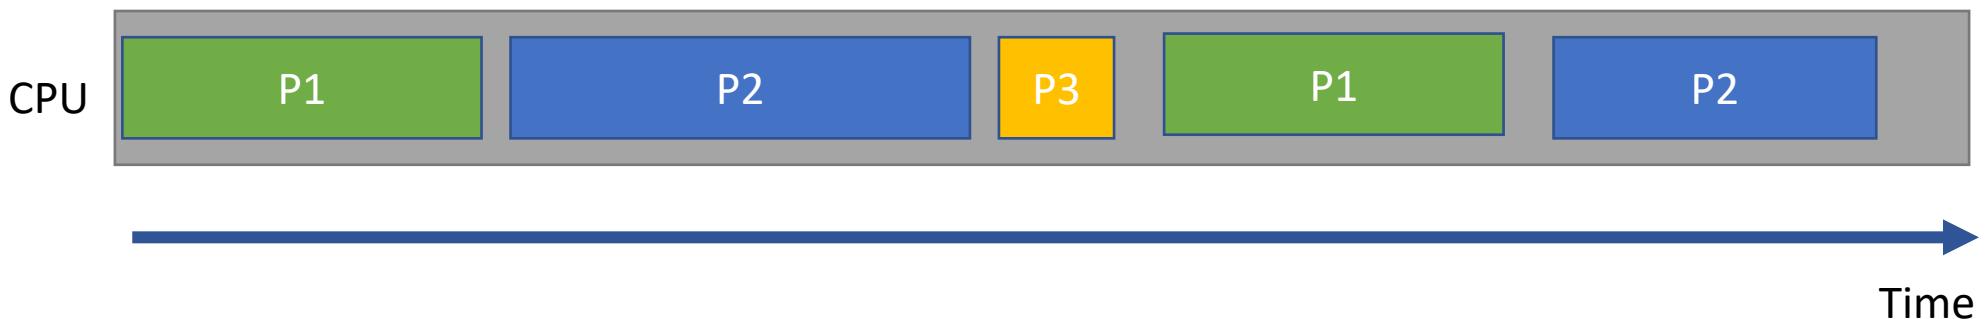
CS 1550

Week 7 – Priority Scheduling with xv6

Teaching Assistant
Maher Khan

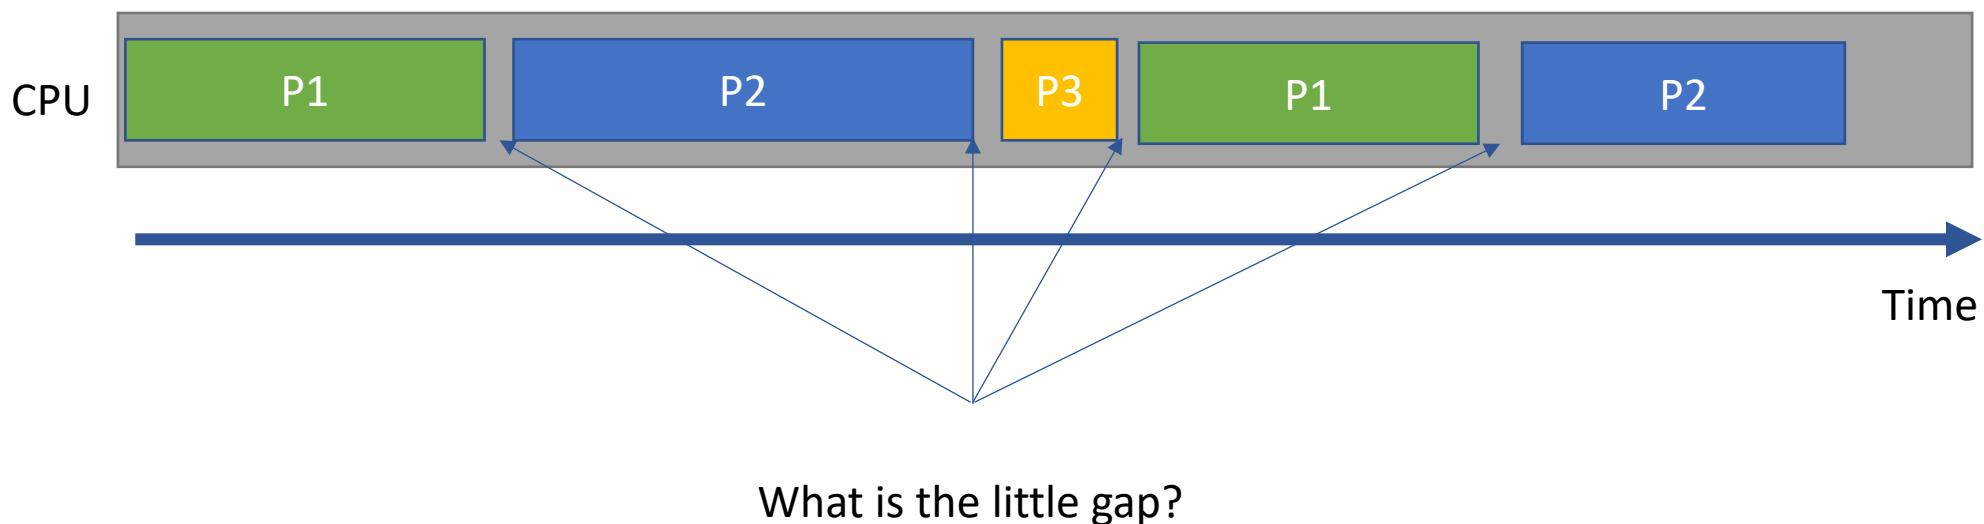
Scheduling of processes

- Switch processes during their execution.
- Currently, processes run in Round Robin inside your XV6



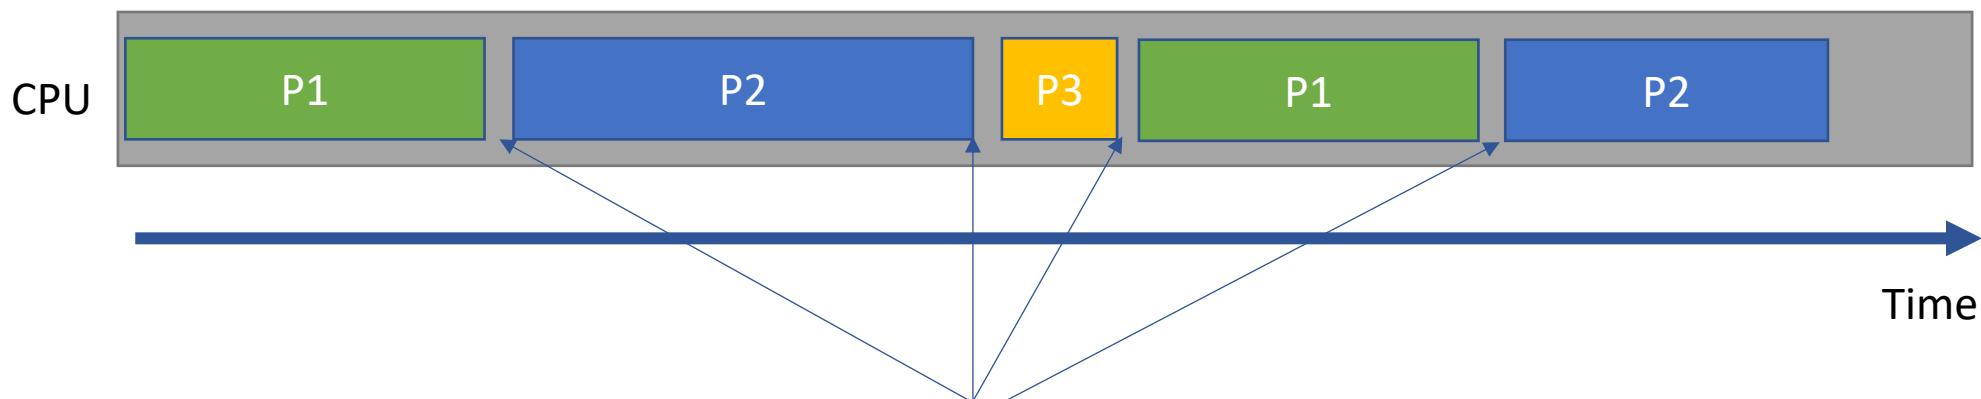
Scheduling of processes

- Switch processes during their execution.
- Currently, processes run in Round Robin inside your XV6



Scheduling of processes

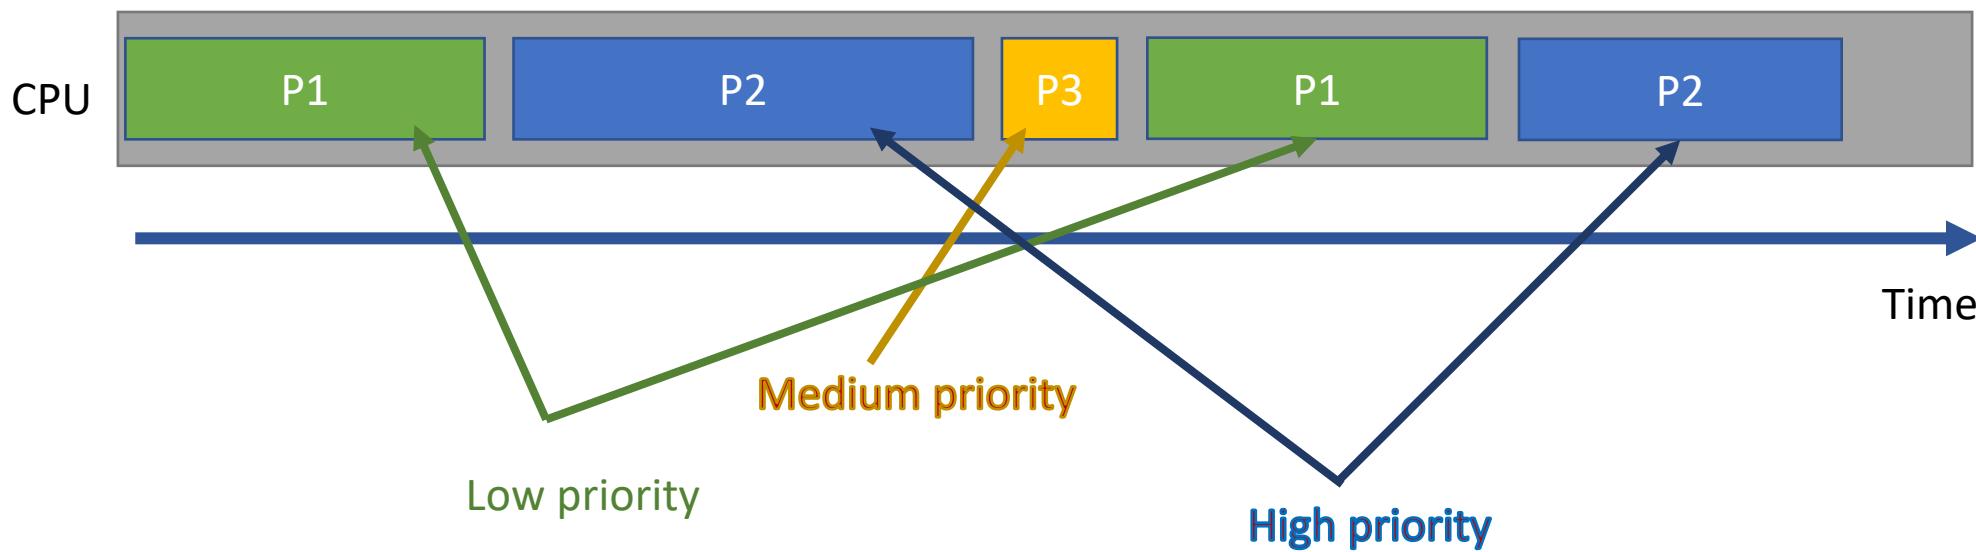
- Switch processes during their execution.
- Currently, processes run in Round Robin inside your XV6



What is the little gap?
The OS **Scheduler** is called when a timer
interrupt occurs, and the kernel calls yield

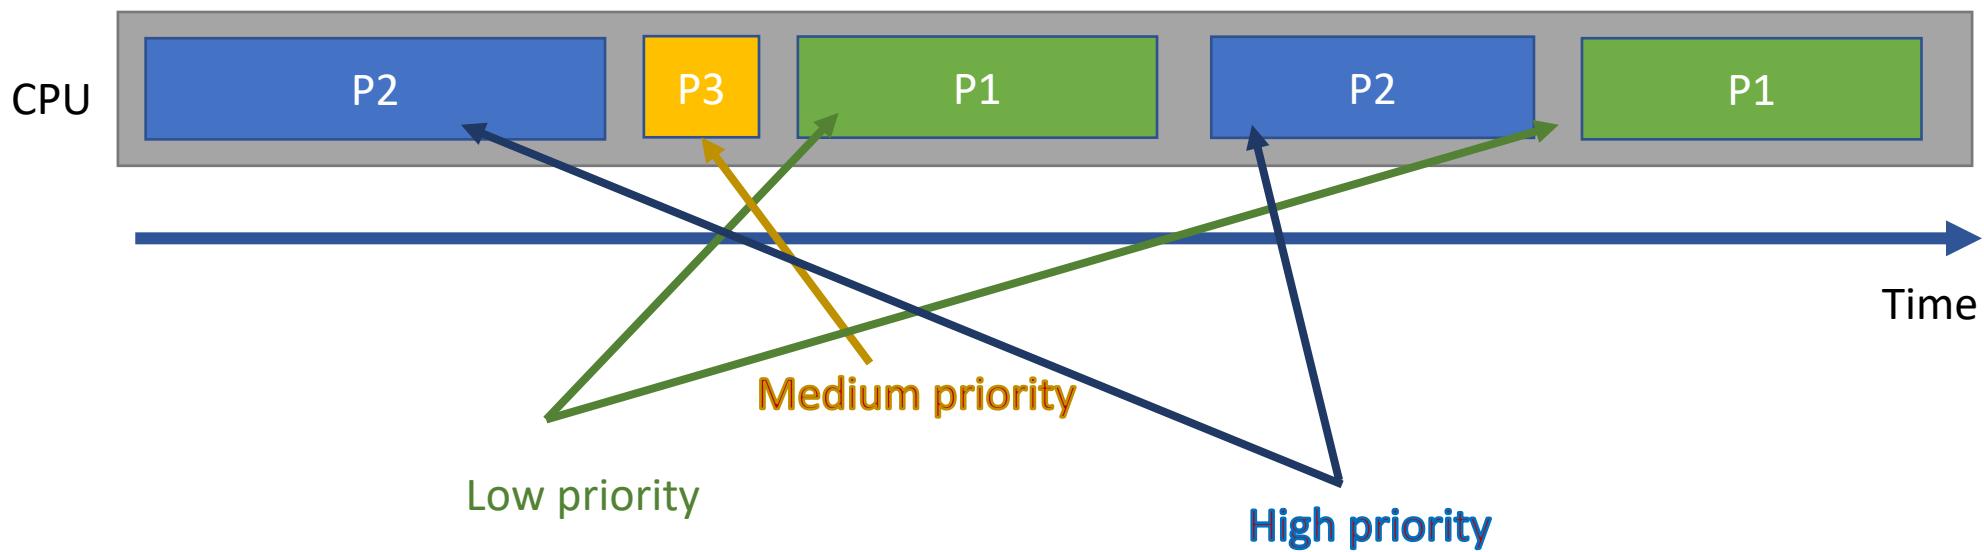
Priority scheduling of processes

- What if processes have different priorities?



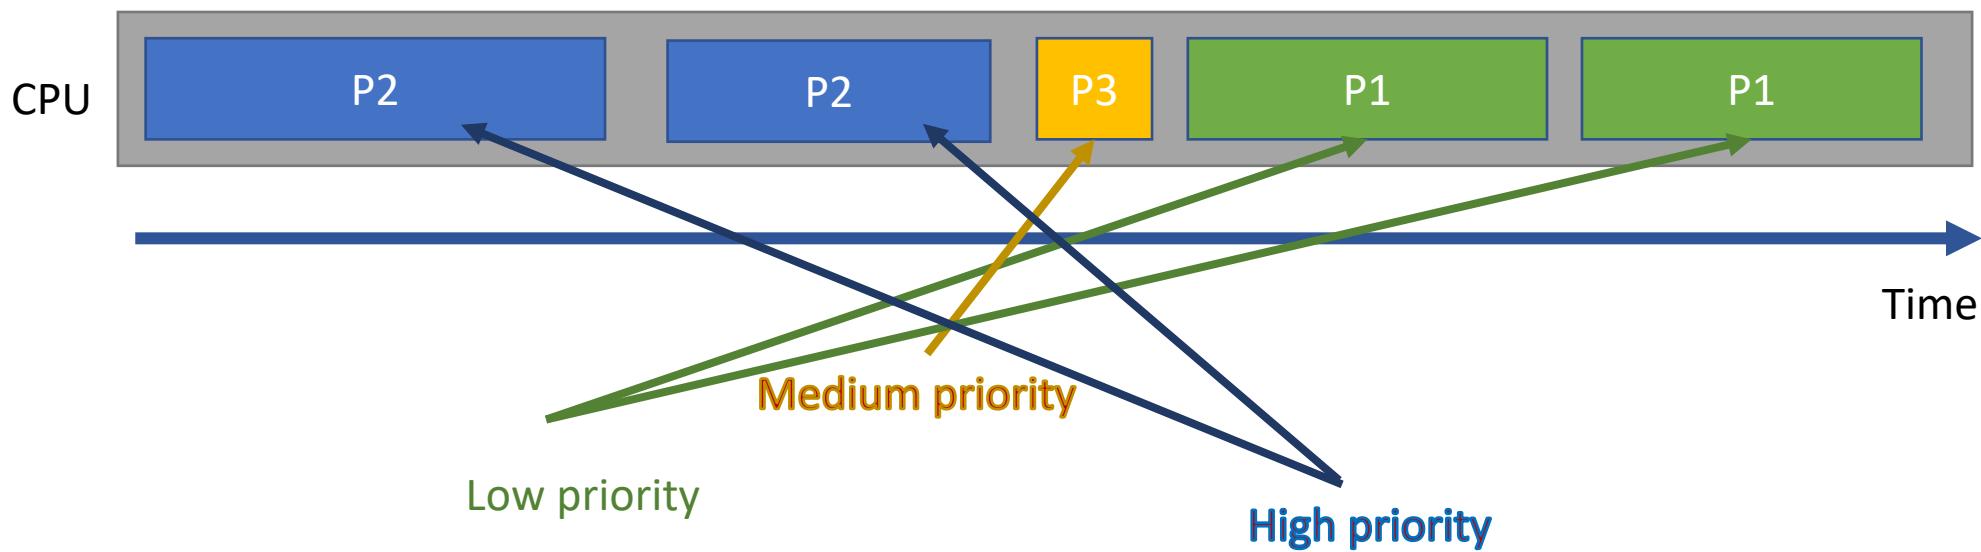
Priority scheduling of processes

- Is this a better arrangement of processes? Can we do better?



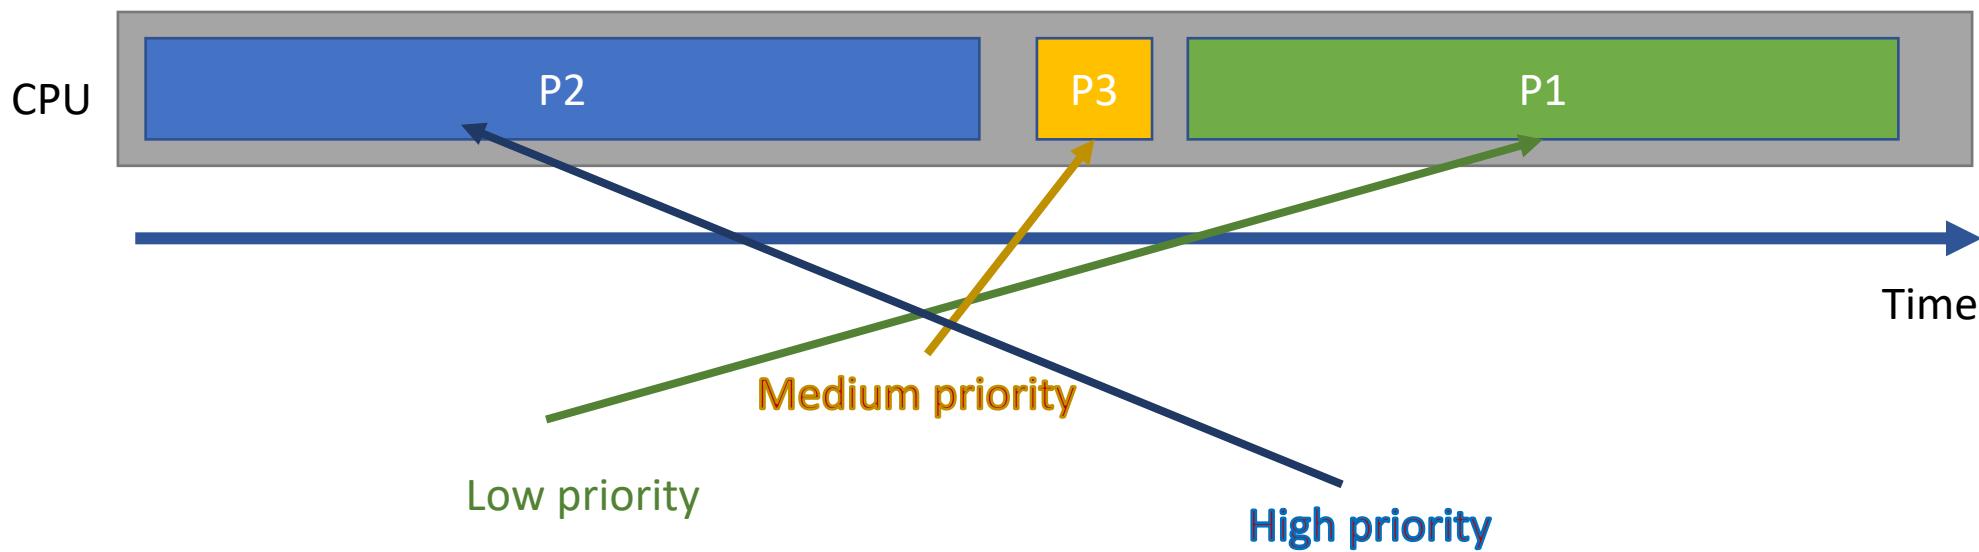
Priority scheduling of processes

- Let all the higher priority processes finish before moving to lower priority ones



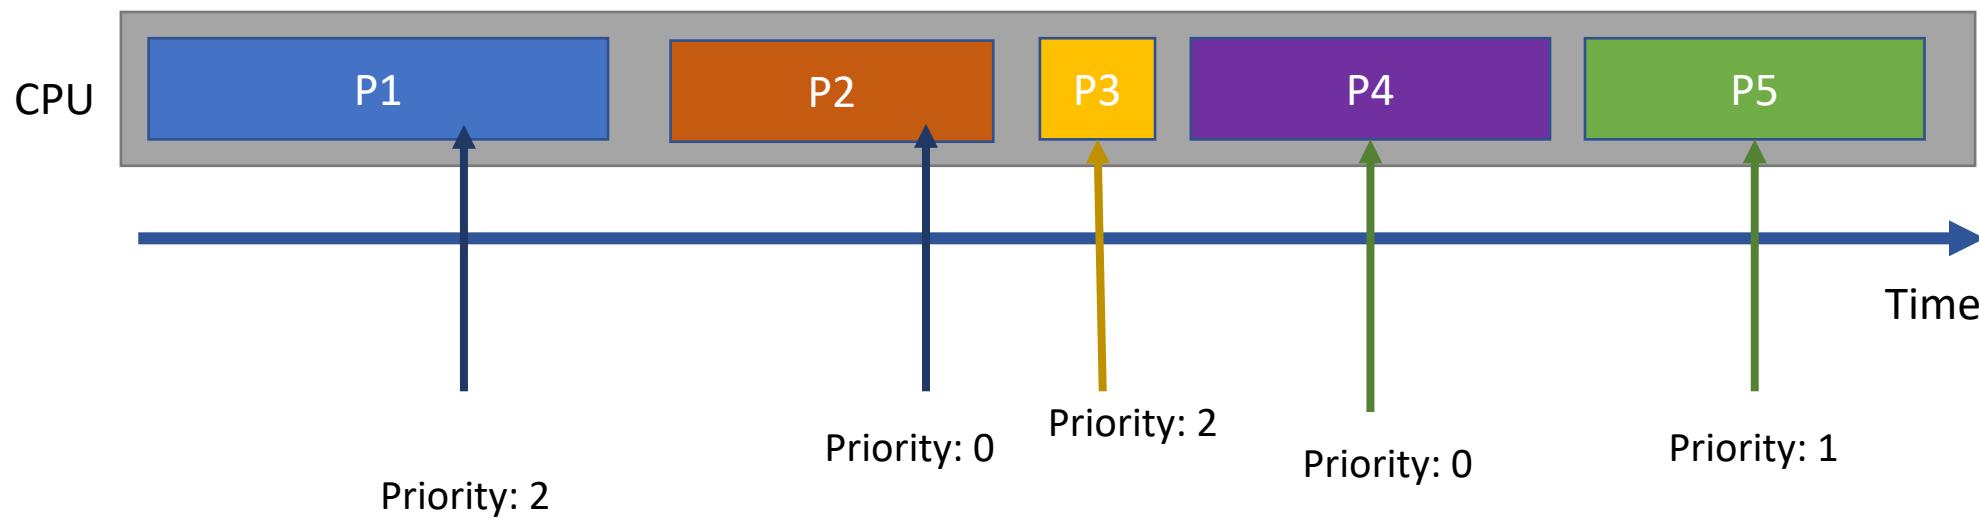
Priority scheduling of processes

- Even better: Don't yield if the current process is the **only one** of its priority!
- This is the bonus part of your lab!



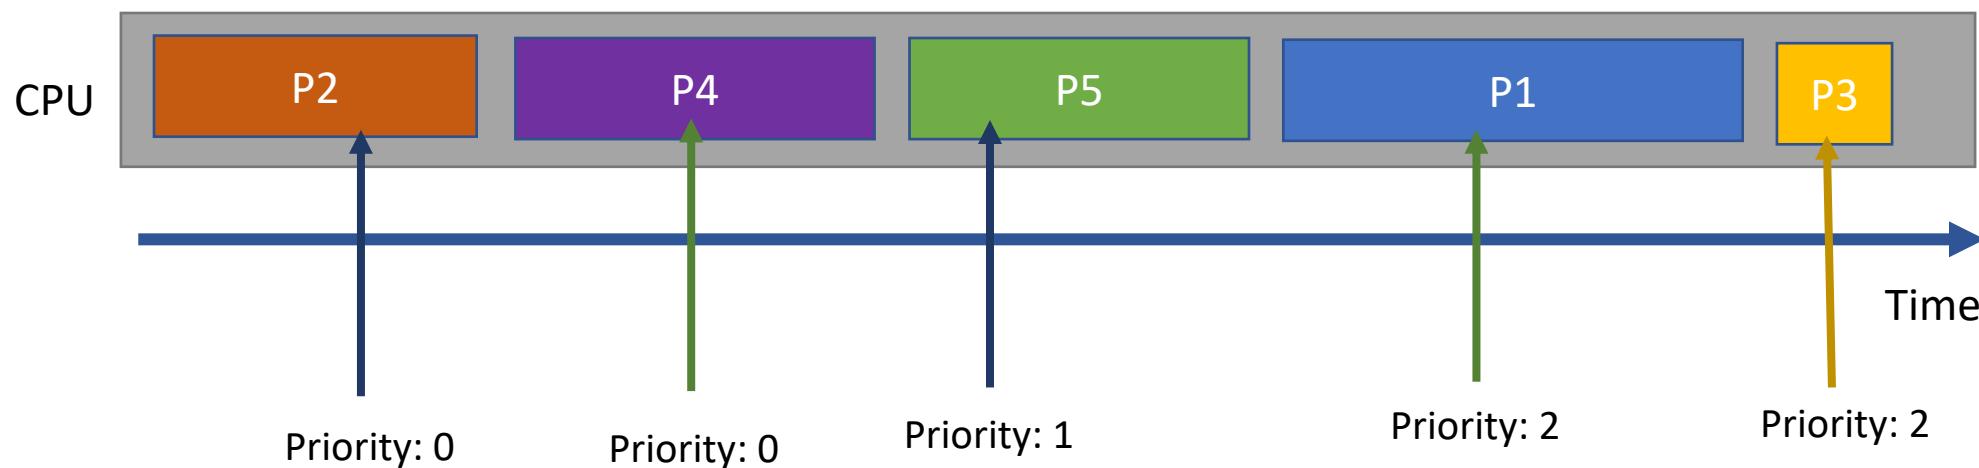
Processes with same priorities

- What if different processes have the same priorities?



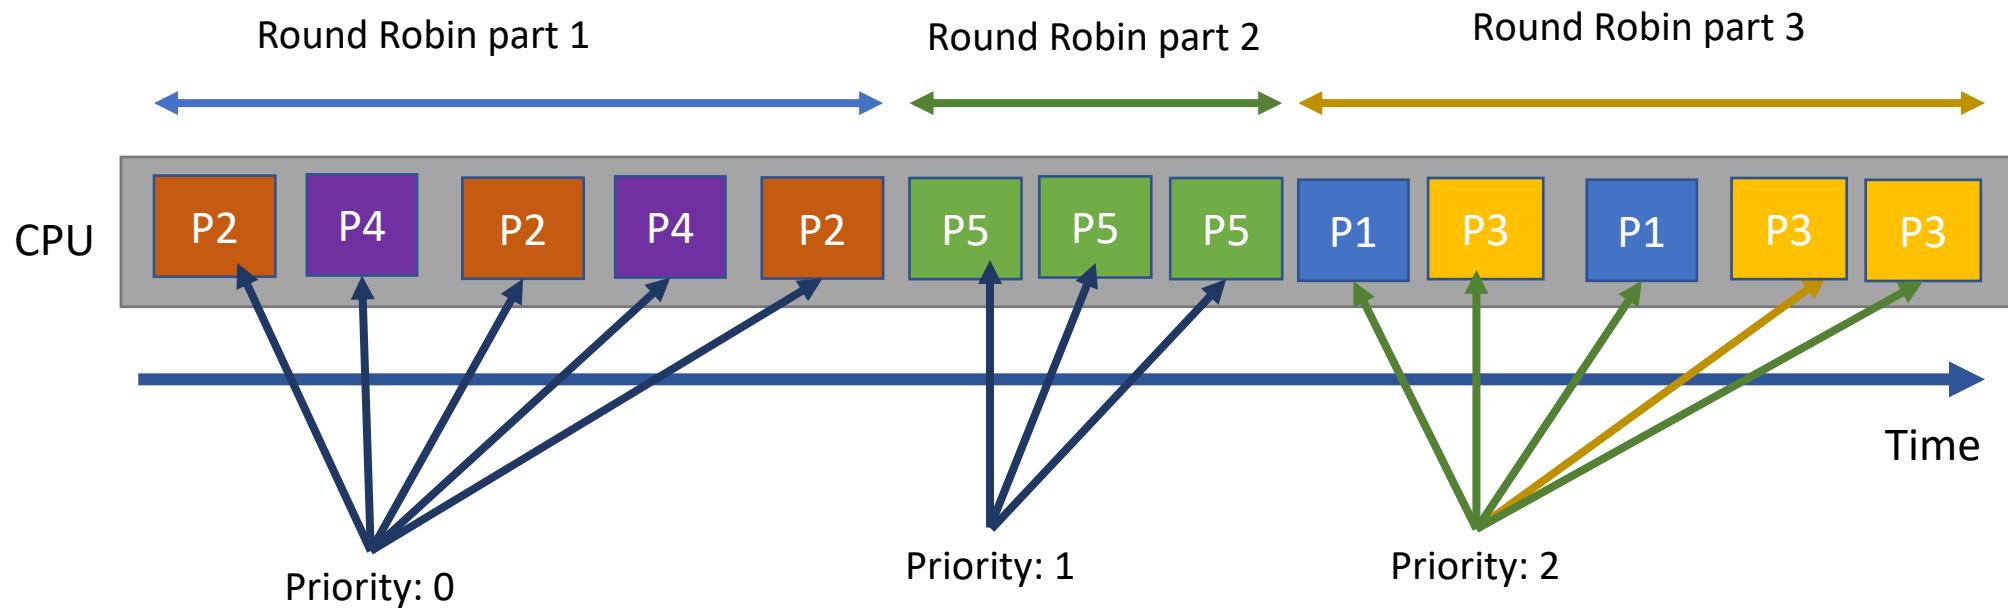
Processes with same priorities

- Group processes with the same priorities together!



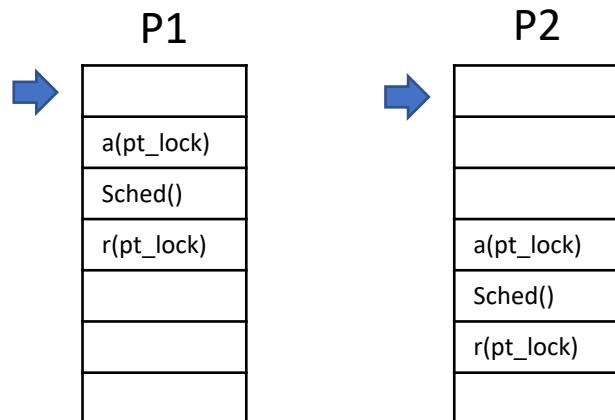
Processes with same priorities

- Processes with same priority should run in Round Robin!



The scheduler function

- Per-CPU process scheduler.
- Each CPU calls scheduler() after setting itself up.
- Scheduler never returns. It loops, doing:
 - choose a process to run
 - call swtch() to start running that process
 - eventually that process transfers control
 - via swtch back to the scheduler.



In proc.c

```
void
scheduler(void)
{
    struct proc *p;
    struct cpu *c = mycpu();
    c->proc = 0;|           Frans Kaashoek (2 years ago) · Eliminate code 1

    for(;;){
        // Enable interrupts on this processor.
        sti();

        // Loop over process table looking for process to run.
        acquire(&ptable.lock);
        for(p = ptable.proc; p < &ptable.proc[NPROC]; p++){
            if(p->state != RUNNABLE)
                continue;

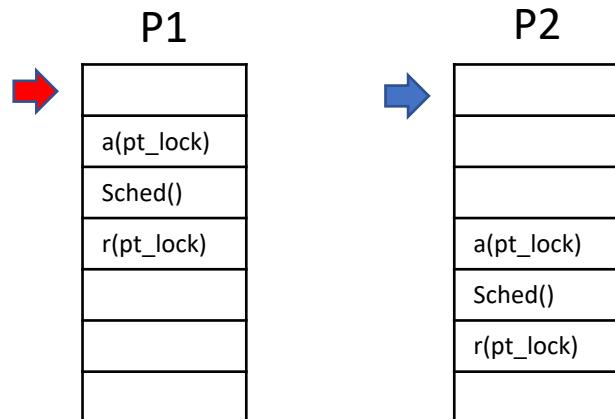
            // Switch to chosen process. It is the process's job
            // to release ptable.lock and then reacquire it
            // before jumping back to us.
            c->proc = p;
            switchuvm(p);
            p->state = RUNNING;

            swtch(&(c->scheduler), p->context);
            switchkvm();

            // Process is done running for now.
            // It should have changed its p->state before coming back.
            c->proc = 0;
        }
        release(&ptable.lock);
    }
}
```

The scheduler function

- Per-CPU process scheduler.
- Each CPU calls scheduler() after setting itself up.
- Scheduler never returns. It loops, doing:
 - choose a process to run
 - call swtch() to start running that process
 - eventually that process transfers control
 - via swtch back to the scheduler.



In proc.c

```
void
scheduler(void)
{
    struct proc *p;
    struct cpu *c = mycpu();
    c->proc = 0;|           Frans Kaashoek (2 years ago) · Eliminate code 1

    for(;;){
        // Enable interrupts on this processor.
        sti();

        // Loop over process table looking for process to run.
        acquire(&ptable.lock);
        for(p = ptable.proc; p < &ptable.proc[NPROC]; p++){
            if(p->state != RUNNABLE)
                continue;

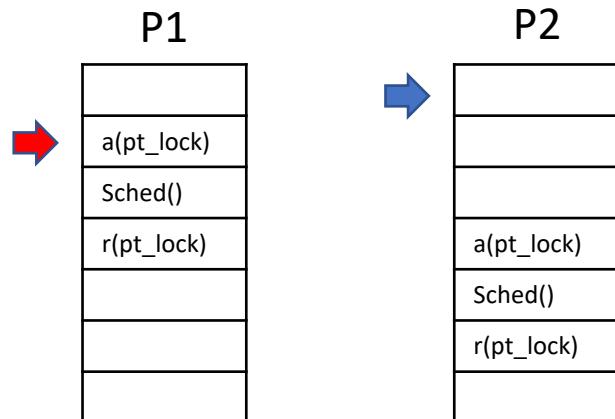
            // Switch to chosen process. It is the process's job
            // to release ptable.lock and then reacquire it
            // before jumping back to us.
            c->proc = p;
            switchuvm(p);
            p->state = RUNNING;

            swtch(&(c->scheduler), p->context);
            switchkvm();

            // Process is done running for now.
            // It should have changed its p->state before coming back.
            c->proc = 0;
        }
        release(&ptable.lock);
    }
}
```

The scheduler function

- Per-CPU process scheduler.
- Each CPU calls scheduler() after setting itself up.
- Scheduler never returns. It loops, doing:
 - choose a process to run
 - call swtch() to start running that process
 - eventually that process transfers control
 - via swtch back to the scheduler.



In proc.c

```
void
scheduler(void)
{
    struct proc *p;
    struct cpu *c = mycpu();
    c->proc = 0;|           Frans Kaashoek (2 years ago) · Eliminate code 1

    for(;;){
        // Enable interrupts on this processor.
        sti();

        // Loop over process table looking for process to run.
        acquire(&ptable.lock);
        for(p = ptable.proc; p < &ptable.proc[NPROC]; p++){
            if(p->state != RUNNABLE)
                continue;

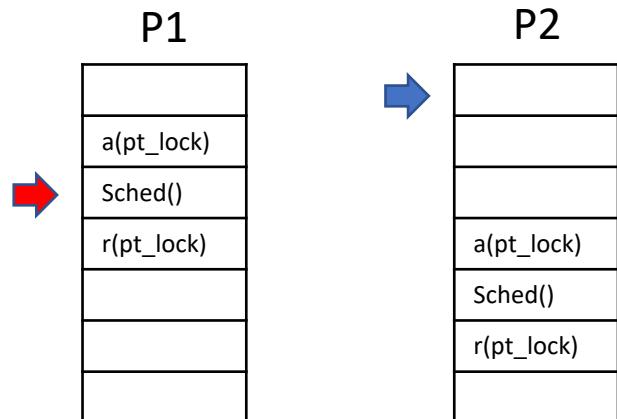
            // Switch to chosen process. It is the process's job
            // to release ptable.lock and then reacquire it
            // before jumping back to us.
            c->proc = p;
            switchuvm(p);
            p->state = RUNNING;

            swtch(&(c->scheduler), p->context);
            switchkvm();

            // Process is done running for now.
            // It should have changed its p->state before coming back.
            c->proc = 0;
        }
        release(&ptable.lock);
    }
}
```

The scheduler function

- Per-CPU process scheduler.
- Each CPU calls scheduler() after setting itself up.
- Scheduler never returns. It loops, doing:
 - choose a process to run
 - call swtch() to start running that process
 - eventually that process transfers control
 - via swtch back to the scheduler.



In proc.c

```
void
scheduler(void)
{
    struct proc *p;
    struct cpu *c = mycpu();
    c->proc = 0;|           Frans Kaashoek (2 years ago) · Eliminate code 1

    for(;;){
        // Enable interrupts on this processor.
        sti();

        // Loop over process table looking for process to run.
        acquire(&ptable.lock);
        for(p = ptable.proc; p < &ptable.proc[NPROC]; p++){
            if(p->state != RUNNABLE)
                continue;

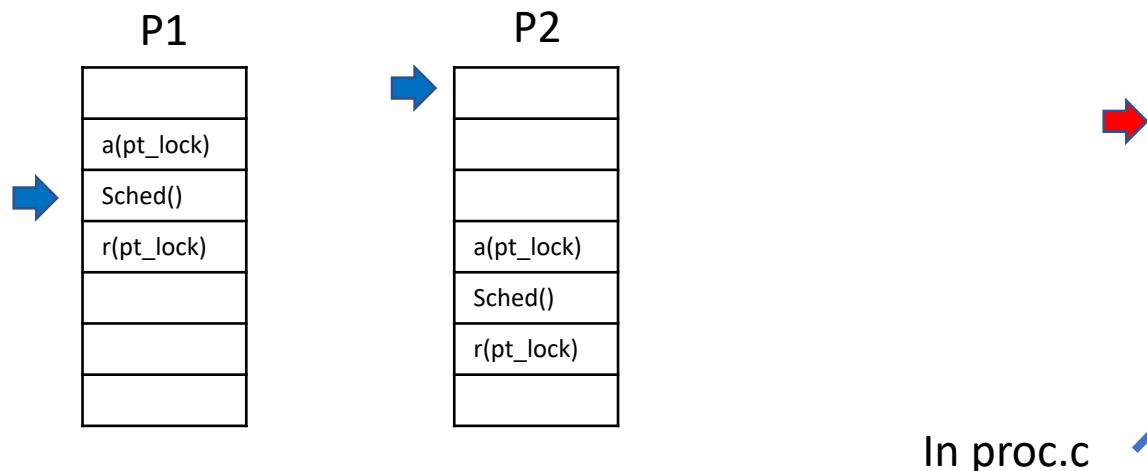
            // Switch to chosen process. It is the process's job
            // to release ptable.lock and then reacquire it
            // before jumping back to us.
            c->proc = p;
            switchuvm(p);
            p->state = RUNNING;

            swtch(&(c->scheduler), p->context);
            switchkvm();

            // Process is done running for now.
            // It should have changed its p->state before coming back.
            c->proc = 0;
        }
        release(&ptable.lock);
    }
}
```

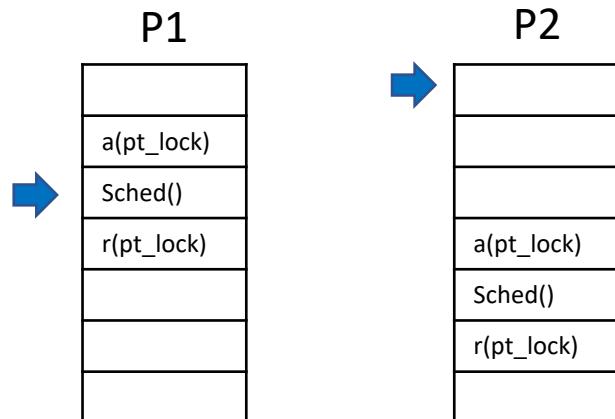
The scheduler function

- Per-CPU process scheduler.
- Each CPU calls scheduler() after setting itself up.
- Scheduler never returns. It loops, doing:
 - choose a process to run
 - call swtch() to start running that process
 - eventually that process transfers control
 - via swtch back to the scheduler.



The scheduler function

- Per-CPU process scheduler.
- Each CPU calls scheduler() after setting itself up.
- Scheduler never returns. It loops, doing:
 - choose a process to run
 - call swtch() to start running that process
 - eventually that process transfers control
 - via swtch back to the scheduler.



In proc.c

```
void
scheduler(void)
{
    struct proc *p;
    struct cpu *c = mycpu();
    c->proc = 0;|           Frans Kaashoek (2 years ago) · Eliminate code 1

    for(;;){
        // Enable interrupts on this processor.
        sti();

        // Loop over process table looking for process to run.
        acquire(&ptable.lock);
        for(p = ptable.proc; p < &ptable.proc[NPROC]; p++){
            if(p->state != RUNNABLE)
                continue;

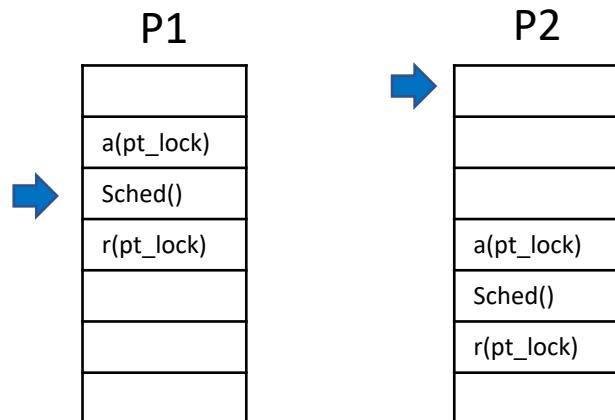
            // Switch to chosen process. It is the process's job
            // to release ptable.lock and then reacquire it
            // before jumping back to us.
            c->proc = p;
            switchuvm(p);
            p->state = RUNNING;

            swtch(&(c->scheduler), p->context);
            switchkvm();

            // Process is done running for now.
            // It should have changed its p->state before coming back.
            c->proc = 0;
        }
        release(&ptable.lock);
    }
}
```

The scheduler function

- Per-CPU process scheduler.
- Each CPU calls scheduler() after setting itself up.
- Scheduler never returns. It loops, doing:
 - choose a process to run
 - call swtch() to start running that process
 - eventually that process transfers control
 - via swtch back to the scheduler.



In proc.c

```
void
scheduler(void)
{
    struct proc *p;
    struct cpu *c = mycpu();
    c->proc = 0;|           Frans Kaashoek (2 years ago) · Eliminate code 1

    for(;;){
        // Enable interrupts on this processor.
        sti();

        // Loop over process table looking for process to run.
        acquire(&ptable.lock);
        for(p = ptable.proc; p < &ptable.proc[NPROC]; p++){
            if(p->state != RUNNABLE)
                continue;

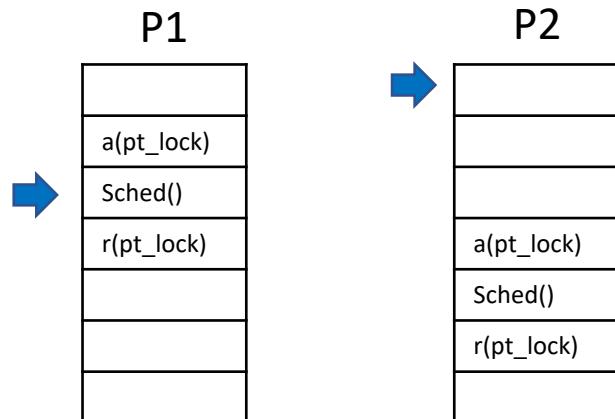
            // Switch to chosen process. It is the process's job
            // to release ptable.lock and then reacquire it
            // before jumping back to us.
            c->proc = p;
            switchuvm(p);
            p->state = RUNNING;

            swtch(&(c->scheduler), p->context);
            switchkvm();

            // Process is done running for now.
            // It should have changed its p->state before coming back.
            c->proc = 0;
        }
        release(&ptable.lock);
    }
}
```

The scheduler function

- Per-CPU process scheduler.
- Each CPU calls scheduler() after setting itself up.
- Scheduler never returns. It loops, doing:
 - choose a process to run
 - call swtch() to start running that process
 - eventually that process transfers control
 - via swtch back to the scheduler.



In proc.c

```
void
scheduler(void)
{
    struct proc *p;
    struct cpu *c = mycpu();
    c->proc = 0;|           Frans Kaashoek (2 years ago) · Eliminate code 1

    for(;;){
        // Enable interrupts on this processor.
        sti();

        // Loop over process table looking for process to run.
        acquire(&ptable.lock);
        for(p = ptable.proc; p < &ptable.proc[NPROC]; p++){
            if(p->state != RUNNABLE)
                continue;

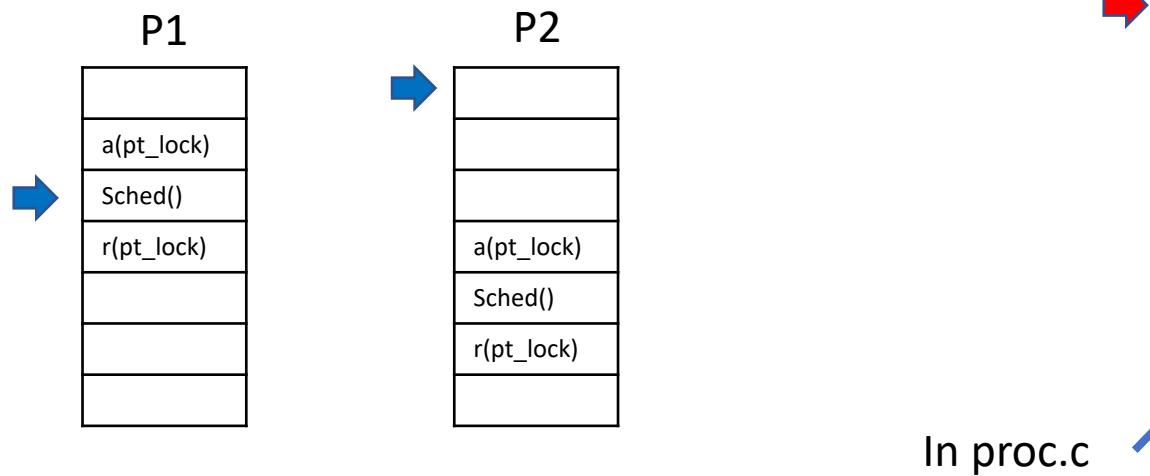
            // Switch to chosen process. It is the process's job
            // to release ptable.lock and then reacquire it
            // before jumping back to us.
            c->proc = p;
            switchuvm(p);
            p->state = RUNNING;

            swtch(&(c->scheduler), p->context);
            switchkvm();

            // Process is done running for now.
            // It should have changed its p->state before coming back.
            c->proc = 0;
        }
        release(&ptable.lock);
    }
}
```

The scheduler function

- Per-CPU process scheduler.
- Each CPU calls scheduler() after setting itself up.
- Scheduler never returns. It loops, doing:
 - choose a process to run
 - call swtch() to start running that process
 - eventually that process transfers control
 - via swtch back to the scheduler.



```
void
scheduler(void)
{
    struct proc *p;
    struct cpu *c = mycpu();
    c->proc = 0;|           Frans Kaashoek (2 years ago) · Eliminate code 1

    for(;;){
        // Enable interrupts on this processor.
        sti();

        // Loop over process table looking for process to run.
        acquire(&ptable.lock);
        for(p = ptable.proc; p < &ptable.proc[NPROC]; p++){
            if(p->state != RUNNABLE)
                continue;

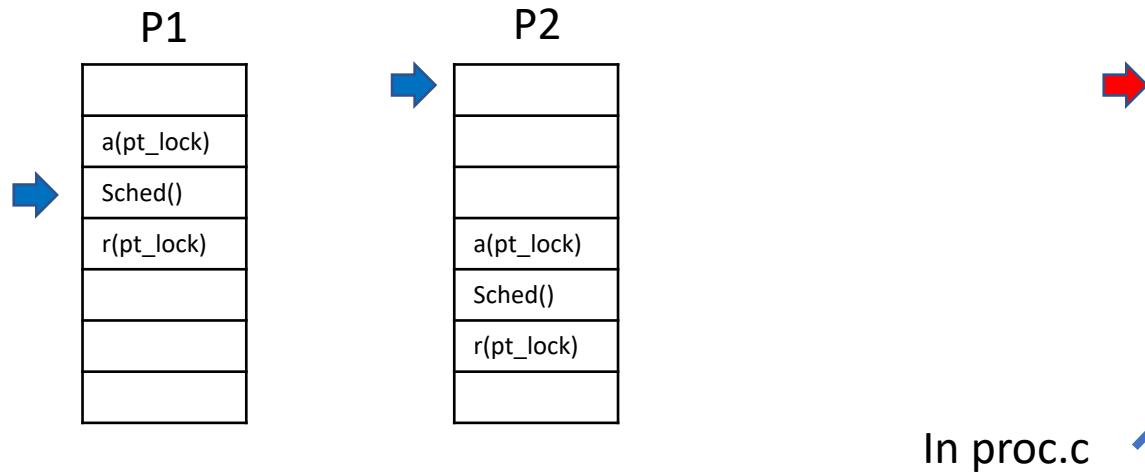
            // Switch to chosen process. It is the process's job
            // to release ptable.lock and then reacquire it
            // before jumping back to us.
            c->proc = p;
            switchuvm(p);
            p->state = RUNNING;

            swtch(&(c->scheduler), p->context);
            switchkvm();

            // Process is done running for now.
            // It should have changed its p->state before coming back.
            c->proc = 0;
        }
        release(&ptable.lock);
    }
}
```

The scheduler function

- Per-CPU process scheduler.
- Each CPU calls scheduler() after setting itself up.
- Scheduler never returns. It loops, doing:
 - choose a process to run
 - call swtch() to start running that process
 - eventually that process transfers control
 - via swtch back to the scheduler.



```
void
scheduler(void)
{
    struct proc *p;
    struct cpu *c = mycpu();
    c->proc = 0;|           Frans Kaashoek (2 years ago) · Eliminate code 1

    for(;;){
        // Enable interrupts on this processor.
        sti();

        // Loop over process table looking for process to run.
        acquire(&ptable.lock);
        for(p = ptable.proc; p < &ptable.proc[NPROC]; p++){
            if(p->state != RUNNABLE)
                continue;

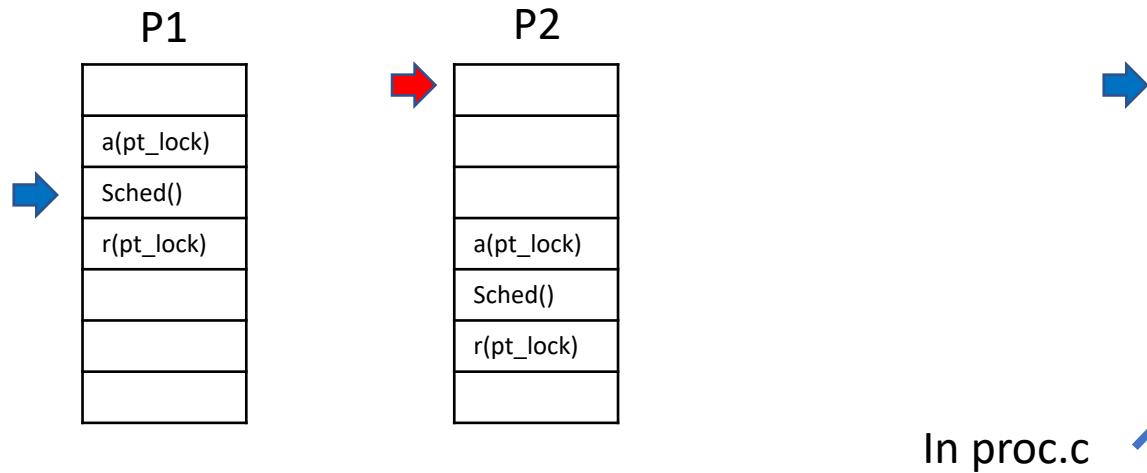
            // Switch to chosen process. It is the process's job
            // to release ptable.lock and then reacquire it
            // before jumping back to us.
            c->proc = p;
            switchuvm(p);
            p->state = RUNNING;

            swtch(&(c->scheduler), p->context);
            switchkvm();

            // Process is done running for now.
            // It should have changed its p->state before coming back.
            c->proc = 0;
        }
        release(&ptable.lock);
    }
}
```

The scheduler function

- Per-CPU process scheduler.
- Each CPU calls scheduler() after setting itself up.
- Scheduler never returns. It loops, doing:
 - choose a process to run
 - call swtch() to start running that process
 - eventually that process transfers control
 - via swtch back to the scheduler.



```
void
scheduler(void)
{
    struct proc *p;
    struct cpu *c = mycpu();
    c->proc = 0;|           Frans Kaashoek (2 years ago) · Eliminate code 1

    for(;;){
        // Enable interrupts on this processor.
        sti();

        // Loop over process table looking for process to run.
        acquire(&ptable.lock);
        for(p = ptable.proc; p < &ptable.proc[NPROC]; p++){
            if(p->state != RUNNABLE)
                continue;

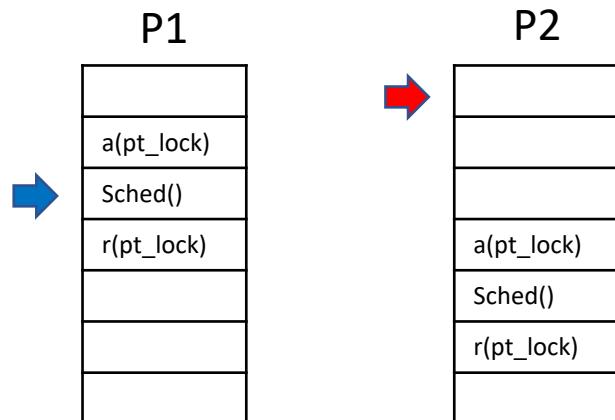
            // Switch to chosen process. It is the process's job
            // to release ptable.lock and then reacquire it
            // before jumping back to us.
            c->proc = p;
            switchuvm(p);
            p->state = RUNNING;

            swtch(&(c->scheduler), p->context);
            switchkvm();

            // Process is done running for now.
            // It should have changed its p->state before coming back.
            c->proc = 0;
        }
        release(&ptable.lock);
    }
}
```

The scheduler function

- Per-CPU process scheduler.
- Each CPU calls scheduler() after setting itself up.
- Scheduler never returns. It loops, doing:
 - choose a process to run
 - call swtch() to start running that process
 - eventually that process transfers control
 - via swtch back to the scheduler.



In proc.c

```
void
scheduler(void)
{
    struct proc *p;
    struct cpu *c = mycpu();
    c->proc = 0;|           Frans Kaashoek (2 years ago) · Eliminate code 1

    for(;;){
        // Enable interrupts on this processor.
        sti();

        // Loop over process table looking for process to run.
        acquire(&ptable.lock);
        for(p = ptable.proc; p < &ptable.proc[NPROC]; p++){
            if(p->state != RUNNABLE)
                continue;

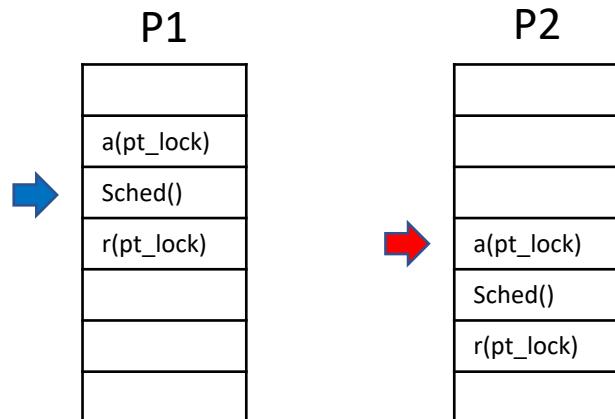
            // Switch to chosen process. It is the process's job
            // to release ptable.lock and then reacquire it
            // before jumping back to us.
            c->proc = p;
            switchuvm(p);
            p->state = RUNNING;

            swtch(&(c->scheduler), p->context);
            switchkvm();

            // Process is done running for now.
            // It should have changed its p->state before coming back.
            c->proc = 0;
        }
        release(&ptable.lock);
    }
}
```

The scheduler function

- Per-CPU process scheduler.
- Each CPU calls scheduler() after setting itself up.
- Scheduler never returns. It loops, doing:
 - choose a process to run
 - call swtch() to start running that process
 - eventually that process transfers control
 - via swtch back to the scheduler.



In proc.c

```
void
scheduler(void)
{
    struct proc *p;
    struct cpu *c = mycpu();
    c->proc = 0;|           Frans Kaashoek (2 years ago) · Eliminate code 1

    for(;;){
        // Enable interrupts on this processor.
        sti();

        // Loop over process table looking for process to run.
        acquire(&ptable.lock);
        for(p = ptable.proc; p < &ptable.proc[NPROC]; p++){
            if(p->state != RUNNABLE)
                continue;

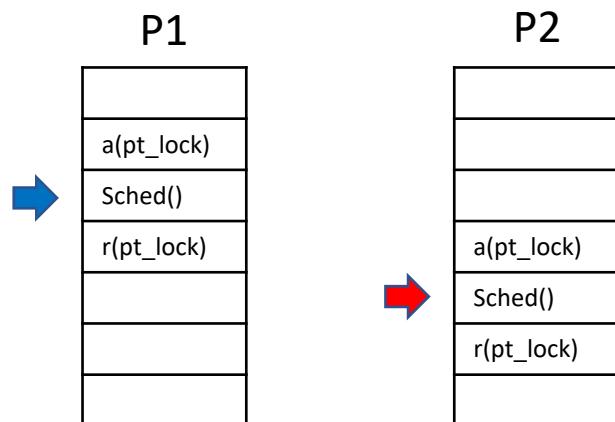
            // Switch to chosen process. It is the process's job
            // to release ptable.lock and then reacquire it
            // before jumping back to us.
            c->proc = p;
            switchuvm(p);
            p->state = RUNNING;

            swtch(&(c->scheduler), p->context);
            switchkvm();

            // Process is done running for now.
            // It should have changed its p->state before coming back.
            c->proc = 0;
        }
        release(&ptable.lock);
    }
}
```

The scheduler function

- Per-CPU process scheduler.
- Each CPU calls scheduler() after setting itself up.
- Scheduler never returns. It loops, doing:
 - choose a process to run
 - call swtch() to start running that process
 - eventually that process transfers control
 - via swtch back to the scheduler.



In proc.c

```
void
scheduler(void)
{
    struct proc *p;
    struct cpu *c = mycpu();
    c->proc = 0;|           Frans Kaashoek (2 years ago) · Eliminate code 1

    for(;;){
        // Enable interrupts on this processor.
        sti();

        // Loop over process table looking for process to run.
        acquire(&ptable.lock);
        for(p = ptable.proc; p < &ptable.proc[NPROC]; p++){
            if(p->state != RUNNABLE)
                continue;

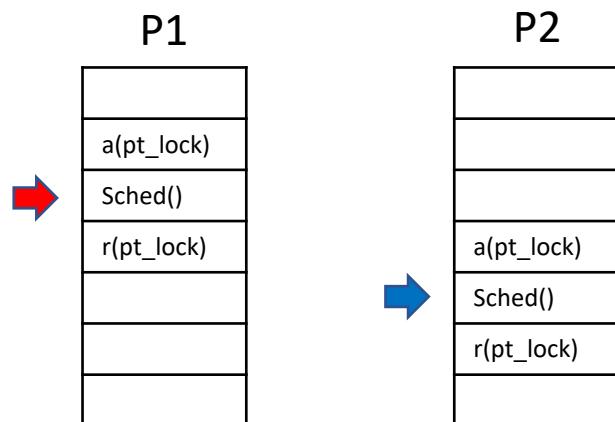
            // Switch to chosen process. It is the process's job
            // to release ptable.lock and then reacquire it
            // before jumping back to us.
            c->proc = p;
            switchuvm(p);
            p->state = RUNNING;

            swtch(&(c->scheduler), p->context);
            switchkvm();

            // Process is done running for now.
            // It should have changed its p->state before coming back.
            c->proc = 0;
        }
        release(&ptable.lock);
    }
}
```

The scheduler function

- Per-CPU process scheduler.
- Each CPU calls scheduler() after setting itself up.
- Scheduler never returns. It loops, doing:
 - choose a process to run
 - call swtch() to start running that process
 - eventually that process transfers control
 - via swtch back to the scheduler.



In proc.c

```
void
scheduler(void)
{
    struct proc *p;
    struct cpu *c = mycpu();
    c->proc = 0;|           Frans Kaashoek (2 years ago) · Eliminate code 1

    for(;;){
        // Enable interrupts on this processor.
        sti();

        // Loop over process table looking for process to run.
        acquire(&ptable.lock);
        for(p = ptable.proc; p < &ptable.proc[NPROC]; p++){
            if(p->state != RUNNABLE)
                continue;

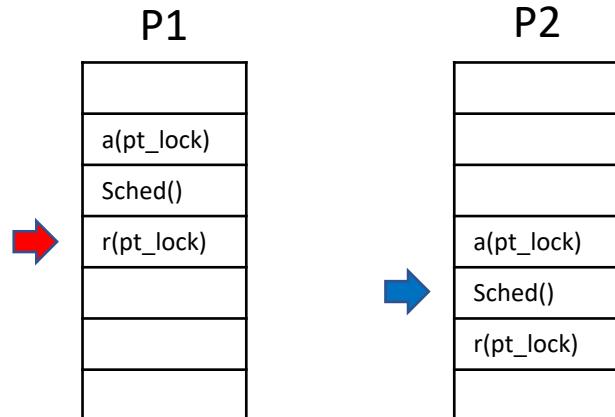
            // Switch to chosen process. It is the process's job
            // to release ptable.lock and then reacquire it
            // before jumping back to us.
            c->proc = p;
            switchuvm(p);
            p->state = RUNNING;

            swtch(&(c->scheduler), p->context);
            switchkvm();

            // Process is done running for now.
            // It should have changed its p->state before coming back.
            c->proc = 0;
        }
        release(&ptable.lock);
    }
}
```

The scheduler function

- Per-CPU process scheduler.
- Each CPU calls scheduler() after setting itself up.
- Scheduler never returns. It loops, doing:
 - choose a process to run
 - call swtch() to start running that process
 - eventually that process transfers control
 - via swtch() back to the scheduler.



In proc.c

```
void
scheduler(void)
{
    struct proc *p;
    struct cpu *c = mycpu();
    c->proc = 0;|           Frans Kaashoek (2 years ago) · Eliminate code 1

    for(;;){
        // Enable interrupts on this processor.
        sti();

        // Loop over process table looking for process to run.
        acquire(&ptable.lock);
        for(p = ptable.proc; p < &ptable.proc[NPROC]; p++){
            if(p->state != RUNNABLE)
                continue;

            // Switch to chosen process. It is the process's job
            // to release ptable.lock and then reacquire it
            // before jumping back to us.
            c->proc = p;
            switchuvm(p);
            p->state = RUNNING;

            swtch(&(c->scheduler), p->context);
            switchkvm();

            // Process is done running for now.
            // It should have changed its p->state before coming back.
            c->proc = 0;
        }
        release(&ptable.lock);
    }
}
```

Yield in trap

- Force process to give up CPU on clock tick.
- IRQ stands for Interrupt Requests
- The timer interrupt occurs after the current process's quantum is over
- The kernel takes over and calls yield() in the trap.c
- Yield() acquires ptable lock and then calls sched()

In trap.c:

```
// ... interrupt were on while ticks ready, needed  
if(myproc() && myproc()->state == RUNNING &&  
| tf->trapno == T_IRQ0+IRQ_TIMER)  
| | yield();
```

In proc.c:

```
// Give up the CPU for one scheduling round.  
void  
yield(void)  
{  
    acquire(&ptable.lock); //DOC: yieldlock  
    myproc()->state = RUNNABLE;  
    sched();  
    release(&ptable.lock);  
}
```

Lab 3 is out!

- Due: 11:59 PM Friday (March 8th)



CS 1550

Week 7 – Priority Scheduling with xv6

Teaching Assistant
Maher Khan